Full job description

We are looking for a skilled **AI/ML Engineer / Data Scientist** to join our team. This role is ideal for someone who enjoys building **real-world AI solutions**, especially leveraging **Generative AI and LLMs**, and can work across the full lifecycle, from experimentation to deployment.

Key Responsibilities

- Design, develop, and deploy **machine learning and AI models** for real-world applications
- Build and implement **Generative AI solutions** (LLMs, chatbots, RAG pipelines, etc.)
- Work with structured and unstructured data for **model training and evaluation**
- Develop scalable backend services/APIs for ML models (using FastAPI/Flask or similar)
- Collaborate with cross-functional teams to translate business problems into AI solutions
- Optimize model performance and ensure production readiness
- Stay updated with the latest advancements in **AI/ML and Generative AI**

Required Skills & Qualifications

- 4–6 years of experience in **Data Science / Machine Learning / AI Engineering**
- Strong programming skills in **Python**
- Hands-on experience with ML libraries such as **Scikit-learn, Pandas, NumPy**
- Experience in building ML models (classification, regression, NLP, etc.)
- Practical exposure to **Generative AI / LLMs**
- Prompt engineering
- RAG (Retrieval-Augmented Generation)
- Chatbot or document-based AI solutions
- Familiarity with frameworks like **LangChain, LlamaIndex, or similar**
- Experience in building and consuming **APIs**
- Strong problem-solving and analytical skills

Good to Have

- Experience with cloud platforms (**AWS / Azure / GCP**)
- Exposure to model deployment tools (Docker, MLflow, etc.)
- Understanding of vector databases (FAISS, Pinecone, etc.)
- Experience working in agile development environments

Job Types: Full-time, Permanent

Benefits:

- Flexible schedule
- Health insurance
- Provident Fund

- Work from home

Application Question(s):

- How many years of hands-on experience do you have with Python ML frameworks (PyTorch, TensorFlow, Scikit-learn)?
- Do you have hands-on experience with Generative AI or LLMs (e.g., chatbots, RAG)?
- Which tools/frameworks have you used for GenAI?

Experience:

- Machine learning: 4 years (Required)

Work Location: Remote